

index

April 30, 2022

1 ARMA Models in statsmodels

1.1 Introduction

In this lesson, you'll use your knowledge of the autoregressive (AR) and moving average (MA) models, along with the `statsmodels` library to model time series data.

1.2 Objectives

You will be able to:

- Fit an AR model using `statsmodels`
- Fit an MA model using `statsmodels`

1.3 Generate a first order AR model

Recall that the AR model has the following formula:

$$Y_t = \mu + \phi * Y_{t-1} + \epsilon_t$$

This means that:

$$Y_1 = \mu + \phi * Y_0 + \epsilon_1$$

$$Y_2 = \mu + \phi * (\text{mean-centered version of } Y_1) + \epsilon_2$$

and so on.

Let's assume a mean-zero white noise with a standard deviation of 2. We'll also create a daily datetime index ranging from January 2017 until the end of March 2018.

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

np.random.seed(11225)

# Create a series with the specified dates
dates = pd.date_range('2017-01-01', '2018-03-31')
```

We will generate a first order AR model with $\phi = 0.7$, $\mu = 5$, and $Y_0 = 8$.

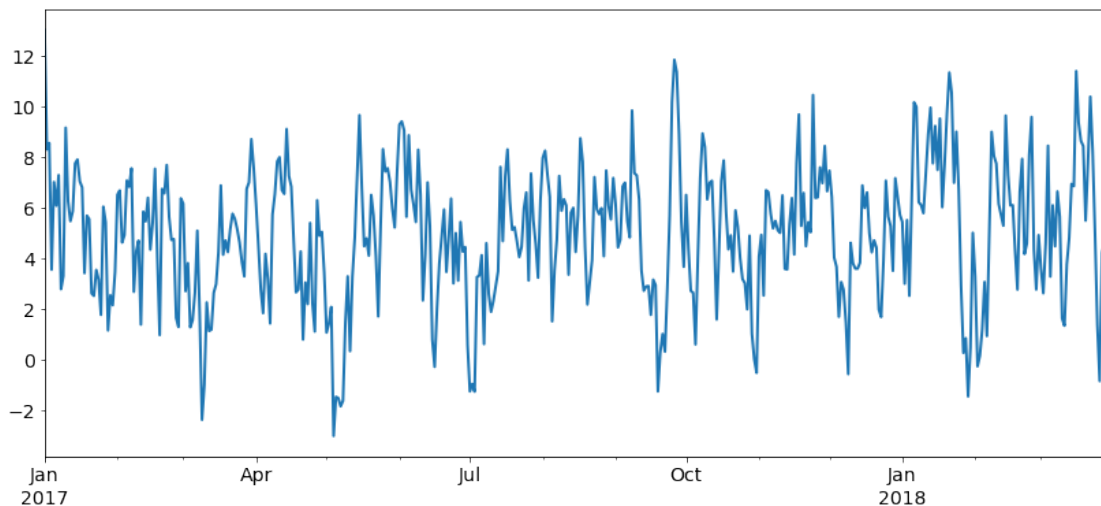
```
[2]: error = np.random.normal(0, 2, len(dates))
     y_0 = 8
     mu = 5
     phi = 0.7
```

```
[3]: TS = [None] * len(dates)
     y = Y_0
     for i, row in enumerate(dates):
         TS[i] = mu + y * phi + error[i]
         y = TS[i] - mu
```

Let's plot the time series to verify:

```
[4]: series = pd.Series(TS, index=dates)

     series.plot(figsize=(14,6), linewidth=2, fontsize=14);
```



1.4 Look at the ACF and PACF of the model

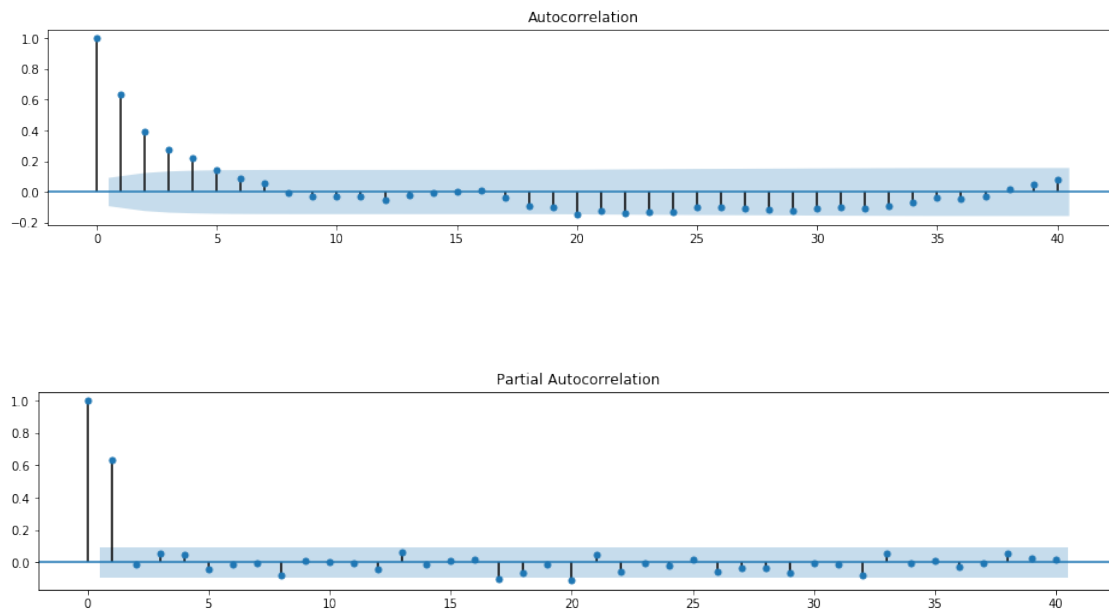
Although we can use `pandas` to plot the ACF, we highly recommended that you use the `statsmodels` variant instead.

```
[6]: from statsmodels.graphics.tsaplots import plot_pacf
     from statsmodels.graphics.tsaplots import plot_acf

     fig, ax = plt.subplots(figsize=(16,3))
     plot_acf(series, ax=ax, lags=40);

     fig, ax = plt.subplots(figsize=(16,3))
```

```
plot_pacf(series, ax=ax, lags=40);
```



1.5 Check the model with ARMA in statsmodels

`statsmodels` also has a tool that fits ARMA models to time series. The only thing we have to do is provide the number of orders for AR and MA. Have a look at the code below, and the output of the code.

The `ARMA()` function requires two arguments: the first is the time series to which the model is fit, and the second is the `order` in the form (p,q) – where `p` refers to the order of AR and `q` refers to the order of MA. For example, a first order AR model would be represented as $(1,0)$.

```
[9]: # Import ARMA
from statsmodels.tsa.arima_model import ARMA
import statsmodels.api as sm
import warnings

# Instantiate an AR(1) model to the simulated data
mod_arma = ARMA(series, order=(1,0))
```

```
/opt/anaconda3/envs/learn-env/lib/python3.8/site-
packages/statsmodels/tsa/arima_model.py:472: FutureWarning:
statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have
been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the .
between arima and model) and
statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.
```

`statsmodels.tsa.arima.model.ARIMA` makes use of the statespace framework and

is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)

warnings.warn(ARIMA_DEPRECATION_WARN, FutureWarning)
```

Once you have instantiated the AR(1) model, you can call the `.fit()` method to fit the model to the data.

```
[10]: # Fit the model to data
      res_arma = mod_arma.fit()
```

Similar to other models, you can then call the `.summary()` method to print the information of the model.

```
[11]: # Print out summary information on the fit
      print(res_arma.summary())
```

```

                        ARMA Model Results
=====
Dep. Variable:          y      No. Observations:          455
Model:                  ARMA(1, 0)  Log Likelihood          -968.698
Method:                  css-mle   S.D. of innovations          2.033
Date:                   Sat, 30 Apr 2022  AIC              1943.395
Time:                   01:59:38    BIC              1955.756
Sample:                 01-01-2017  HQIC              1948.265
                        - 03-31-2018
=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
const         4.9664      0.269     18.444      0.000        4.439        5.494
ar.L1.y        0.6474      0.036     17.880      0.000        0.576        0.718
                        Roots
=====
              Real      Imaginary      Modulus      Frequency
-----
AR.1          1.5446      +0.0000j        1.5446        0.0000
=====
```

Make sure that the output for the ϕ parameter and μ is as you'd expect. You can use the `.params` attribute to check these values.

```
[12]: # Print out the estimate for the constant and for theta
print(res_arma.params)
```

```
const      4.966377
ar.L1.y     0.647429
dtype: float64
```

1.6 Generate a first order MA model

Recall that the MA model has the following formula:

$$Y_t = \mu + \epsilon_t + \theta * \epsilon_{t-1}$$

This means that:

$$Y_1 = \mu + \epsilon_1 + \theta * \epsilon_0$$

$$Y_2 = \mu + \epsilon_2 + \theta * \epsilon_1$$

and so on.

Assume a mean-zero white noise with a standard deviation of 4. We'll also generate a daily datetime index ranging from April 2015 until the end of August 2015.

We will generate a first order MA model with $\theta = 0.9$ and $\mu = 7$.

```
[13]: np.random.seed(1234)

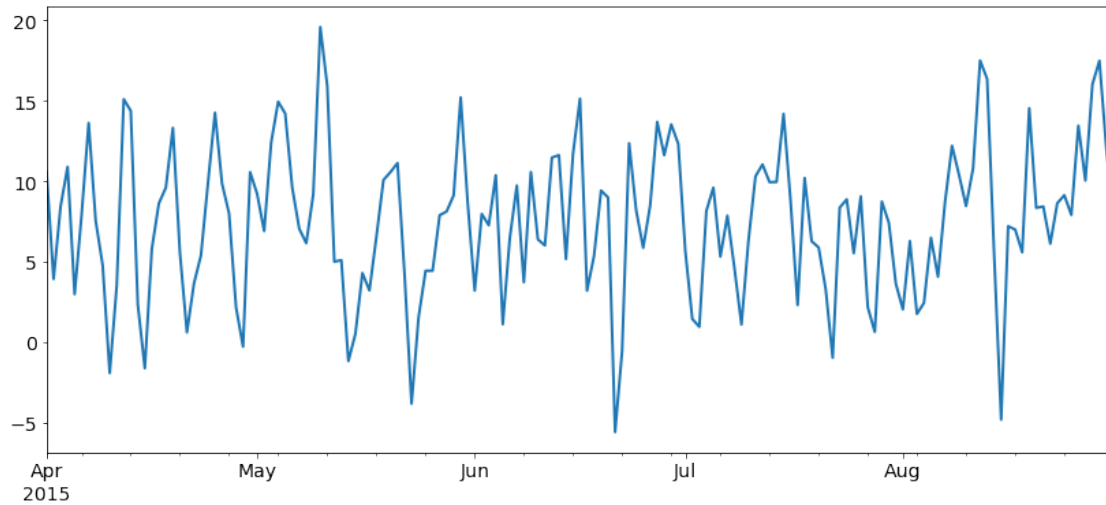
# Create a series with the specified dates
dates = pd.date_range('2015-04-01', '2015-08-31')

error = np.random.normal(0, 4, len(dates))
mu = 7
theta = 0.9

TS = [None] * len(dates)
error_prev = error[0]
for i, row in enumerate(dates):
    TS[i] = mu + theta * error_prev + error[i]
    error_prev = error[i]
```

```
[14]: series = pd.Series(TS, index=dates)

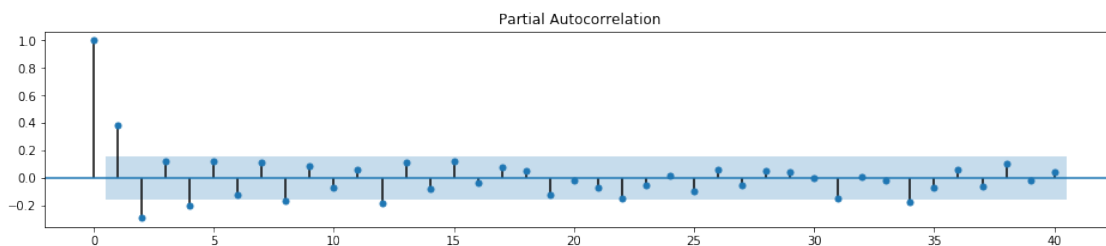
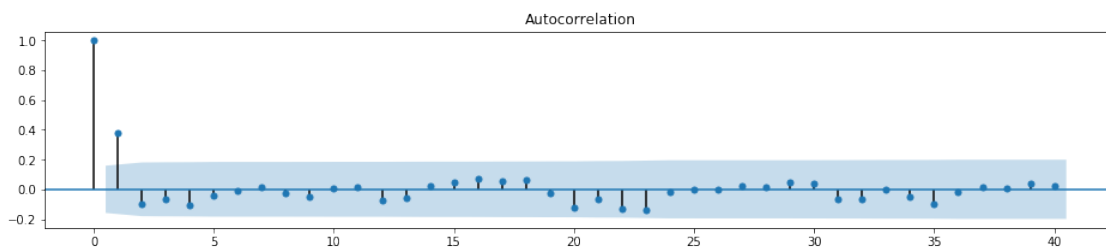
series.plot(figsize=(14,6), linewidth=2, fontsize=14);
```



1.7 Look at the ACF and PACF of the model

```
[15]: fig, ax = plt.subplots(figsize=(16,3))
      plot_acf(series, ax=ax, lags=40);

      fig, ax = plt.subplots(figsize=(16,3))
      plot_pacf(series, ax=ax, lags=40);
```



1.8 Check the model with ARMA in statsmodels

Let's fit an MA model to verify the parameters are estimated correctly. The first order MA model would be represented as (0,1).

```
[16]: # Instantiate and fit an MA(1) model to the simulated data
mod_arma = ARMA(series, order=(0,1))
res_arma = mod_arma.fit()

# Print out summary information on the fit
print(res_arma.summary())
```

```

                        ARMA Model Results
=====
Dep. Variable:          y      No. Observations:          153
Model:                  ARMA(0, 1)  Log Likelihood        -426.378
Method:                  css-mle    S.D. of innovations      3.909
Date:                   Sat, 30 Apr 2022    AIC                858.757
Time:                   02:02:02    BIC                867.848
Sample:                 04-01-2015    HQIC               862.450
                        - 08-31-2015
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const         7.5373         0.590      12.776      0.000         6.381         8.694
ma.L1.y        0.8727         0.051      17.165      0.000         0.773         0.972
                        Roots
=====
              Real      Imaginary      Modulus      Frequency
-----
MA.1         -1.1459         +0.0000j         1.1459         0.5000
=====
```

```
/opt/anaconda3/envs/learn-env/lib/python3.8/site-
packages/statsmodels/tsa/arma_model.py:472: FutureWarning:
statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have
been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the .
between arma and model) and
statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.
```

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARMA',
                        FutureWarning)
```

```
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',  
                        FutureWarning)
```

```
warnings.warn(ARIMA_DEPRECATION_WARN, FutureWarning)
```

```
[17]: # Print out the estimate for the constant and for theta  
      print(res_arma.params)
```

```
const      7.537294  
ma.L1.y     0.872683  
dtype: float64
```

1.9 Summary

Great job! In this lesson, you saw how you can use the AR and MA models using the `ARMA()` function from `statsmodels` by specifying the order in the form of `(p,q)`, where at least one of `p` or `q` was zero depending on the kind of model fit. You can use `ARMA()` to fit a combined ARMA model as well – which you will do in the next lab!